

# Nearest neighbor methods

## Lecture 10

David Sontag  
New York University


Slides adapted from Vibhav Gogate, Carlos Guestrin,  
Mehryar Mohri, & Luke Zettlemoyer

# Nearest Neighbor Algorithm

- Learning Algorithm:
  - Store training examples
- Prediction Algorithm:
  - To classify a new example  $\mathbf{x}$  by finding the training example  $(\mathbf{x}^i, y^i)$  that is *nearest* to  $\mathbf{x}$
  - Guess the class  $y = y^i$

# K-Nearest Neighbor Methods

- To classify a new input vector  $x$ , examine the  $k$ -closest training data points to  $x$  and assign the object to the most frequently occurring class




common values for  $k$ : 3, 5

# Decision Boundaries



- The nearest neighbor algorithm does not explicitly compute decision boundaries. However, the decision boundaries form a subset of the Voronoi diagram for the training data.

*1-NN Decision Surface*



- The more examples that are stored, the more complex the decision boundaries can become

# Example results for k-NN



[Figures from Hastie and Tibshirani, Chapter 13]

# Nearest Neighbor

## When to Consider

- Instance map to points in  $R^n$
- Less than 20 attributes per instance
- Lots of training data

## Advantages

- Training is very fast
- Learn complex target functions
- Do not lose information

## Disadvantages

- Slow at query time
- Easily fooled by irrelevant attributes

# Issues

- Distance measure
  - Most common: Euclidean
- Choosing  $k$ 
  - Increasing  $k$  reduces variance, increases bias
- For high-dimensional space, problem that the nearest neighbor may not be very close at all!
- Memory-based technique. Must make a pass through the data for each classification. This can be prohibitive for large data sets.

## Distance

- Notation: object with  $p$  measurements

$$\mathbf{x}^i = (x_1^i, x_2^i, \dots, x_p^i)$$

- Most common distance metric is *Euclidean* distance:

$$d_E(\mathbf{x}^i, \mathbf{x}^j) = \left( \sum_{k=1}^p (x_k^i - x_k^j)^2 \right)^{\frac{1}{2}}$$

- ED makes sense when different measurements are commensurate; each is variable measured in the same units.
- If the measurements are different, say length and weight, it is not clear.



# Standardization

When variables are not commensurate, we can standardize them by dividing by the sample standard deviation. This makes them all equally important.

The estimate for the standard deviation of  $x_k$  :

$$\hat{\sigma}_k = \left( \frac{1}{n} \sum_{i=1}^n (x_k^i - \bar{x}_k)^2 \right)^{\frac{1}{2}}$$

where  $\bar{x}_k$  is the sample mean:

$$\bar{x}_k = \frac{1}{n} \sum_{i=1}^n x_k^i$$

## Weighted Euclidean distance


Finally, if we have some idea of the relative importance of each variable, we can weight them:

$$d_{WE}(i, j) = \left( \sum_{k=1}^p w_k (x_k^i - x_k^j)^2 \right)^{\frac{1}{2}}$$


# The Curse of Dimensionality

- Nearest neighbor breaks down in high-dimensional spaces because the “neighborhood” becomes very large.
- Suppose we have 5000 points uniformly distributed in the unit hypercube and we want to apply the 5-nearest neighbor algorithm.
- Suppose our query point is at the origin.
  - 1D –
    - On a one dimensional line, we must go a distance of  $5/5000 = 0.001$  on average to capture the 5 nearest neighbors
  - 2D –
    - In two dimensions, we must go  $\sqrt{0.001}$  to get a square that contains 0.001 of the volume
  - D –
    - In D dimensions, we must go  $(0.001)^{1/D}$

## K-NN and irrelevant features



## K-NN and irrelevant features




## Nearest neighbor problem

- **Problem:** given sample  $S = ((x_1, y_1), \dots, (x_m, y_m))$ , find the nearest neighbor of test point  $x$ .
- general problem extensively studied in computer science.
- exact vs. approximate algorithms.
- dimensionality  $N$  crucial.
- better algorithms for small intrinsic dimension (e.g., limited doubling dimension).

## Efficient Indexing: N=2

### ■ Algorithm:

- compute Voronoi diagram in  $O(m \log m)$ .
- **point location** data structure to determine NN.
- complexity:  $O(m)$  space,  $O(\log m)$  time.




## Efficient Indexing: $N > 2$

- Voronoi diagram: size in  $O(m^{\lceil N/2 \rceil})$ .
- Linear algorithm (no pre-processing):
  - compute distance  $\|x - x_i\|$  for all  $i \in [1, m]$ .
  - complexity of distance computation:  $\Omega(Nm)$ .
  - no additional space needed.
- Tree-based data structures: pre-processing.
  - often used in applications:  **$k$ -d trees** ( $k$ -dimensional trees).



# Efficient Indexing for $N > 2$ : KD trees



Construction algorithm


- **Algorithm:** for each non-leaf node,
    - choose dimension (e.g., longest of hyperrectangle).
    - choose pivot (median).
    - split node according to (pivot, dimension).
- ➔ balanced tree, binary space partitioning.

# Efficient Indexing for $N > 2$ : KD trees

## ■ Algorithm:

- find region containing  $x$  (starting from root node, move to child node based on node test).
- save region point  $x_0$  as current best.
- move up tree and recursively search regions intersecting hypersphere  $S(x, \|x - x_0\|)$ :
  - update current best if current point is closer.
  - restart search with each intersecting sub-tree.
  - move up tree when no more intersecting sub-tree.

Search algorithm



## KNN Advantages

- Easy to program
- No optimization or training required
- Classification accuracy can be very good; can outperform more complex models